

CSE 31

Computer Organization

Lecture 5 – C Memory Management



Announcement

- ▶ Lab #2 this week
 - Due in one week
- ▶ HW #1
 - From zyBooks
 - Due today
- ▶ Reading assignment
 - Chapter 7 and 8.7 of K&R (C book) to review on C/C++ programming
- ▶ Tutoring from PALS

C structures : Overview

- ▶ A **struct** is a data structure composed from simpler data types.
 - Like a class in Java/C++ but without methods or inheritance.

```
struct point {    /* type definition */  
    int x;  
    int y;  
};
```

As always in C, the argument is passed by “value” – a copy is made.

```
void PrintPoint(struct point p) {  
  
    printf("( %d, %d) ", p.x, p.y);  
}
```

```
struct point p1 = {0, 10}; /* x=0, y=10 */
```

```
PrintPoint(p1);
```

C structures: Pointers to them

- ▶ Usually, more efficient to pass a pointer to the struct.
- ▶ The C arrow operator ($->$) dereferences and extracts a structure field (member) with a single operator.
- ▶ The following are equivalent:

```
struct point *p;  
/* code to assign to pointer */  
printf("x is %d\n", (*p).x);  
printf("x is %d\n", p->x);
```

How big are structs?

- ▶ Recall C operator `sizeof()` which gives size in bytes (of type or variable)
- ▶ How big is `sizeof(p)`?

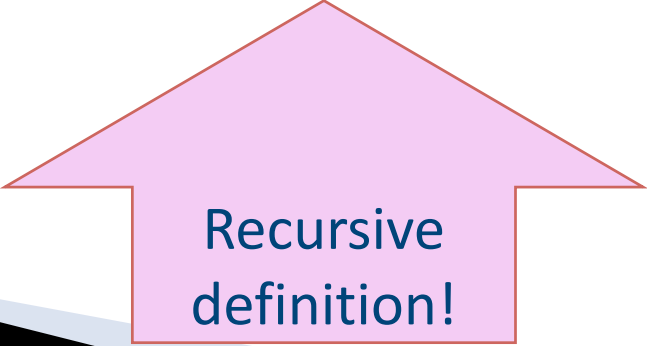
```
struct p {  
    char x;  
    int y;  
};
```

- 5 bytes? 8 bytes?
- Compiler may word align integer `y`

Linked List Example

- ▶ Let's look at an example of using structures, pointers, `malloc()`, and `free()` to implement a **linked list of strings**.

```
/* node structure for linked list */  
struct Node {  
    char *value;  
    struct Node *next;  
};
```



Recursive
definition!

typedef simplifies the code

```
struct Node {  
    char *value;  
    struct Node *next;  
};
```



String value;

```
/* "typedef" means define a new type */  
typedef struct Node NodeStruct;
```

... OR ...

```
typedef struct Node {  
    char *value;  
    struct Node *next;  
} NodeStruct;
```

... THEN

```
typedef NodeStruct *List;  
typedef char *String;
```

```
/* Note similarity! */  
/* To define 2 nodes */
```

```
struct Node {  
    char *value;  
    struct Node *next;  
} node1, node2;
```

Linked List Example

```
/* Add a string to an existing list */
```

```
List cons(String s, List list)
{
    List node = (List) malloc(sizeof(NodeStruct));

    node->value = (String) malloc (strlen(s) + 1);
    strcpy(node->value, s);
    node->next = list;
    return node;
}
```

```
String s1 = "abc", s2 = "cde";
List theList = NULL;
theList = cons(s2, theList);
theList = cons(s1, theList);
    /* or embedded */
theList = cons(s1, cons(s2, NULL));
```

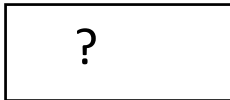

Linked List Example

/ Add a string to an existing list, 2nd call */*

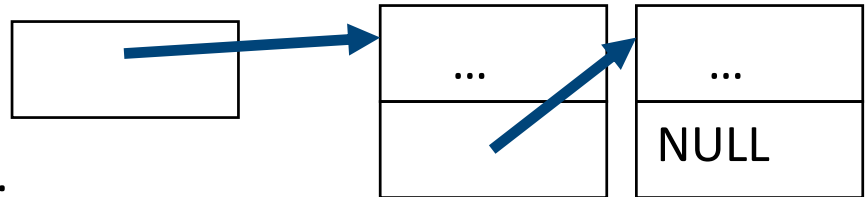
```
List cons(String s, List list)
{
    List node = (List) malloc(sizeof(NodeStruct));

    node->value = (String) malloc (strlen(s) + 1);
    strcpy(node->value, s);
    node->next = list;
    return node;
}
```

node:



list:



s:



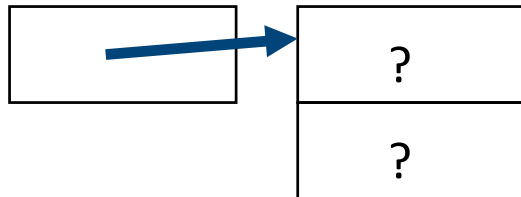
Linked List Example

/ Add a string to an existing list, 2nd call */*

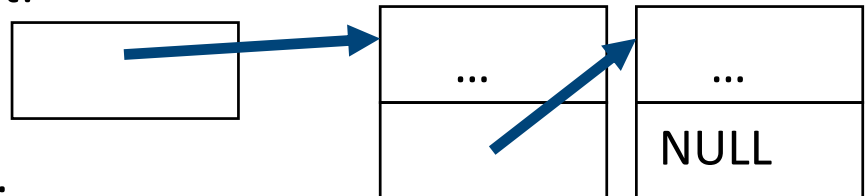
```
List cons(String s, List list)
{
    List node = (List) malloc(sizeof(NodeStruct));

    node->value = (String) malloc (strlen(s) + 1);
    strcpy(node->value, s);
    node->next = list;
    return node;
}
```

node:



list:



s:



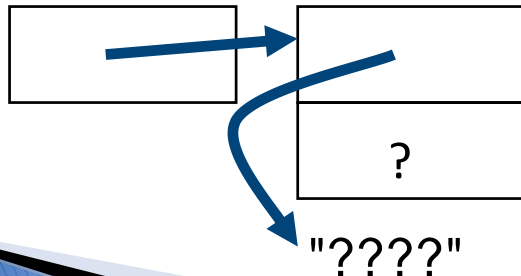
Linked List Example

/ Add a string to an existing list, 2nd call */*

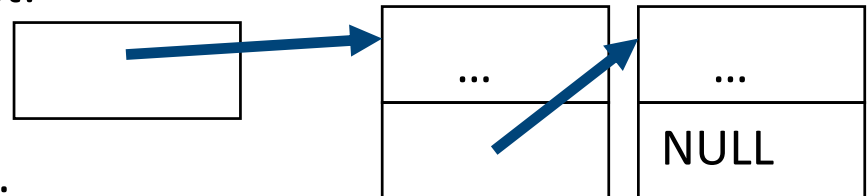
```
List cons(String s, List list)
{
    List node = (List) malloc(sizeof(NodeStruct));

    node->value = (String) malloc (strlen(s) + 1);
    strcpy(node->value, s);
    node->next = list;
    return node;
}
```

node:



list:



s:



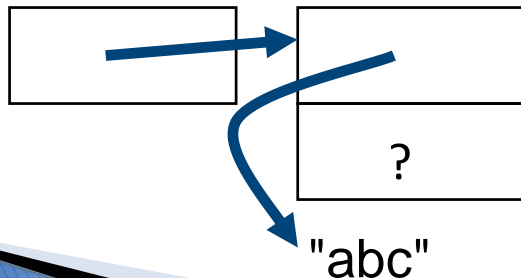
Linked List Example

/ Add a string to an existing list, 2nd call */*

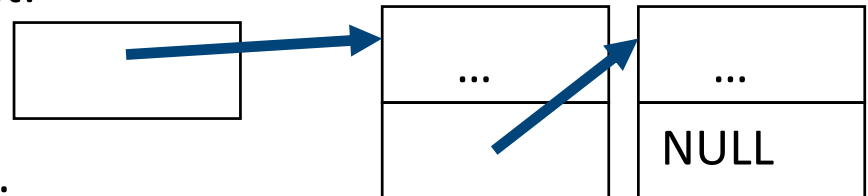
```
List cons(String s, List list)
{
    List node = (List) malloc(sizeof(NodeStruct));

    node->value = (String) malloc (strlen(s) + 1);
    strcpy(node->value, s);
    node->next = list;
    return node;
}
```

node:



list:



s:

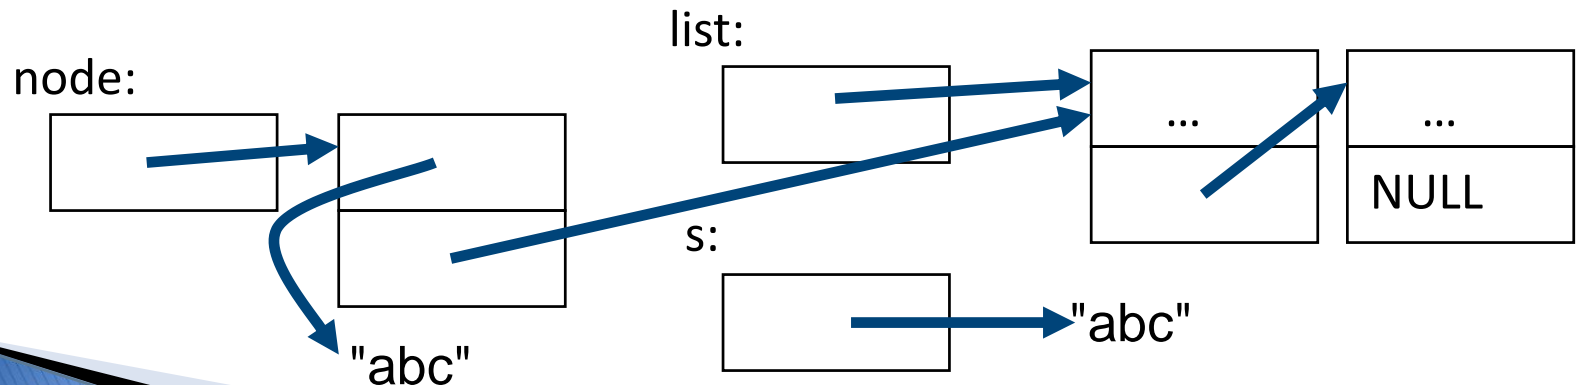


Linked List Example

/ Add a string to an existing list, 2nd call */*

```
List cons(String s, List list)
{
    List node = (List) malloc(sizeof(NodeStruct));

    node->value = (String) malloc (strlen(s) + 1);
    strcpy(node->value, s);
    node->next = list;
    return node;
}
```

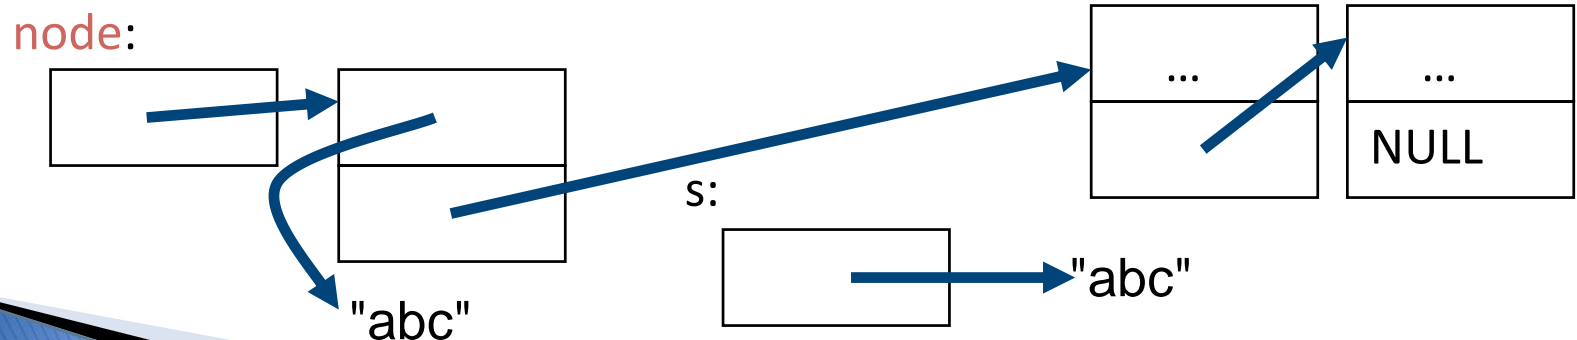


Linked List Example

/ Add a string to an existing list, 2nd call */*

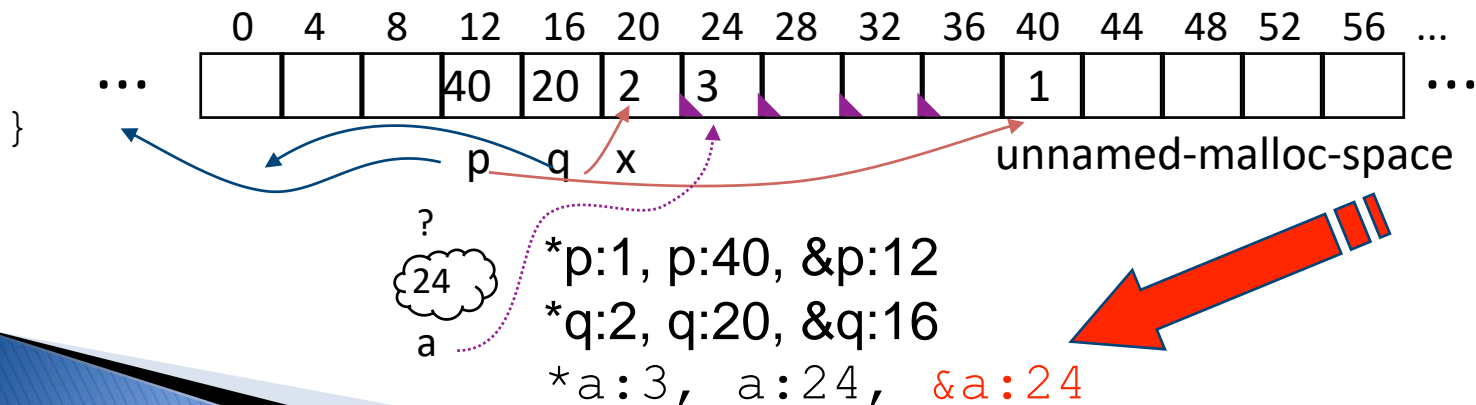
```
List cons(String s, List list)
{
    List node = (List) malloc(sizeof(NodeStruct));

    node->value = (String) malloc (strlen(s) + 1);
    strcpy(node->value, s);
    node->next = list;
    return node;
}
```



Arrays not implemented as you'd think

```
void foo() {  
    int *p, *q, x;  
    int a[4];  
    p = (int *) malloc (sizeof(int));  
    q = &x;  
  
    *p = 1; // p[0] would also work here  
    printf("*p:%u, p:%u, &p:%u\n", *p, p, &p);  
    *q = 2; // q[0] would also work here  
    printf("*q:%u, q:%u, &q:%u\n", *q, q, &q);  
    *a = 3; // a[0] would also work here  
    printf("*a:%u, a:%u, &a:%u\n", *a, a, &a);  
}
```



K&R: "An array name is not a variable"

Don't forget the globals!

▶ Remember:

- Structure declaration does not allocate memory
 - Only when you instantiate it.
- Variable declaration does allocate memory

▶ So far we have talked about several different ways to allocate memory for data:

1. Declaration of a local variable

```
int i; struct Node list; char *string;  
int ar[n];
```

2. “Dynamic” allocation at runtime by calling allocation function (malloc).

```
ptr = (struct Node *) malloc(sizeof(struct  
Node) *n);
```

▶ One more possibility exists...

3. Data declared outside of any procedure/function (i.e., before main).

- Similar to #1 above, but has “global” scope.

Useful in C, but not in Java/C++

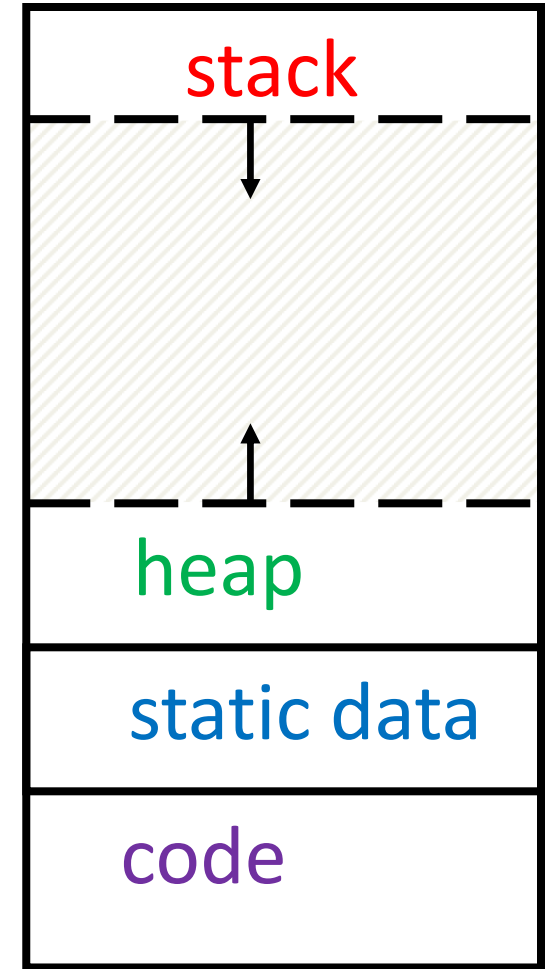
```
int myGlobal;  
main() {  
    ...  
}
```


C Memory Management

- ▶ C has 3 pools of memory (based on the nature of usage)
 - Static storage: global variable storage, basically permanent, entire program run
 - The Stack: local variable storage, parameters, return address (location of “activation records” in Java or “stack frame” in C)
 - The Heap (dynamic malloc storage): data lives until deallocated by programmer
- ▶ C requires knowing where objects are in memory, otherwise things don't work as expected
 - Java hides location of objects

Normal C Memory Management

$\sim FFFF\ FFFF_{hex}$



- ▶ A program's **address space** contains 4 regions:
 - **stack**: local variables, grows downward
 - **heap**: space requested for pointers via `malloc()` ; resizes dynamically, grows upward
 - **static data**: variables declared outside main, does not grow or shrink
 - **code**: loaded when program starts, does not change

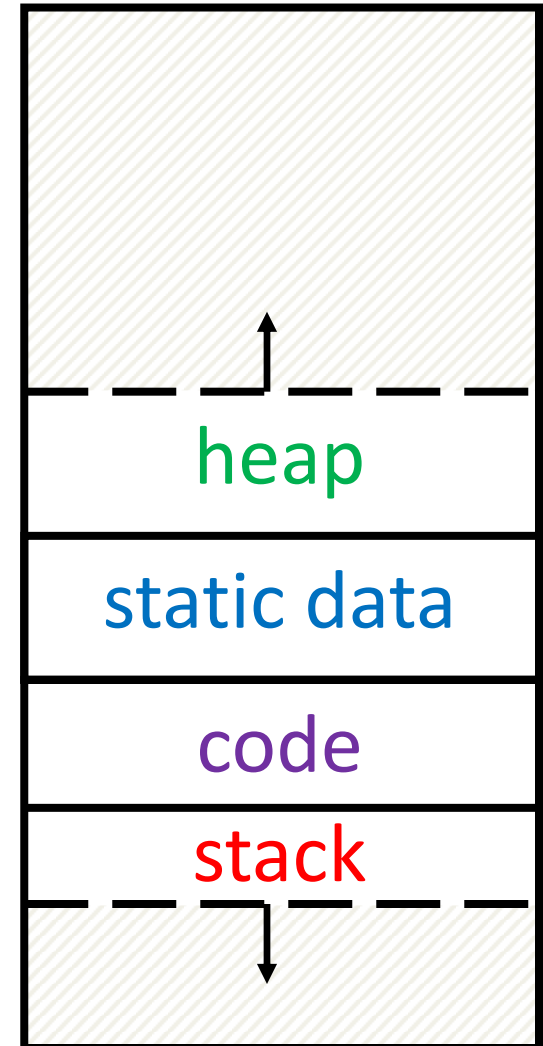
$\sim 0_{hex}$

For now, OS somehow prevents accesses between stack and heap (gray hash lines). Wait for virtual memory

Intel 80x86 C Memory Management

- ▶ A C program's 80x86 address space :
 - **heap**: space requested for pointers via `malloc()`; resizes dynamically, grows upward
 - **static data**: variables declared outside main, does not grow or shrink
 - **code**: loaded when program starts, does not change
 - **stack**: local variables, grows downward

~ 08000000_{hex}



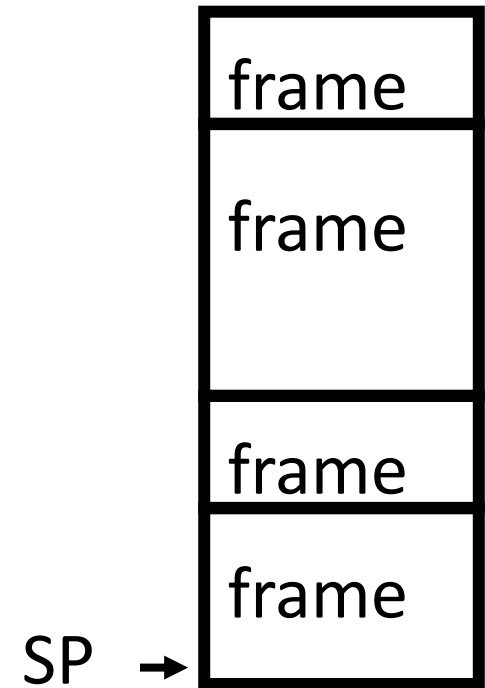
Where are variables allocated?

- ▶ If declared outside of a function
 - allocated in “static” storage
- ▶ If declared inside of a function
 - allocated in the “stack”
 - freed when a function returns.
 - That’s why the scope is within the function
- ▶ Note: `main()` is a function!

```
int myGlobal;  
main() {  
    int myTemp;  
}
```

Stack frames

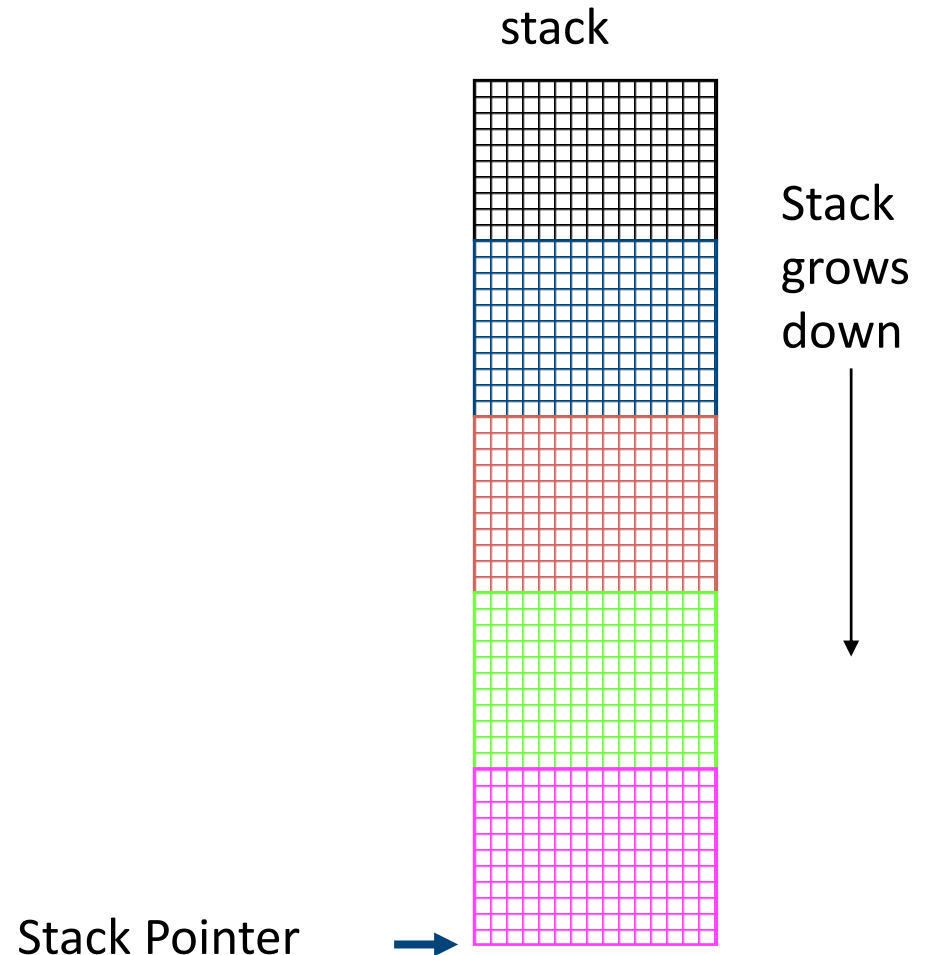
- ▶ Stack frame includes storage for:
 - Return “instruction” address
 - Parameters (input arguments)
 - Space for other local variables
- ▶ Stack frames:
 - contiguous blocks of memory
 - stack pointer tells where top stack frame is
- ▶ When a function ends, stack frame is “**popped off**” the stack; frees memory for future stack frames



Stack

- ▶ Last In, First Out (LIFO) data structure

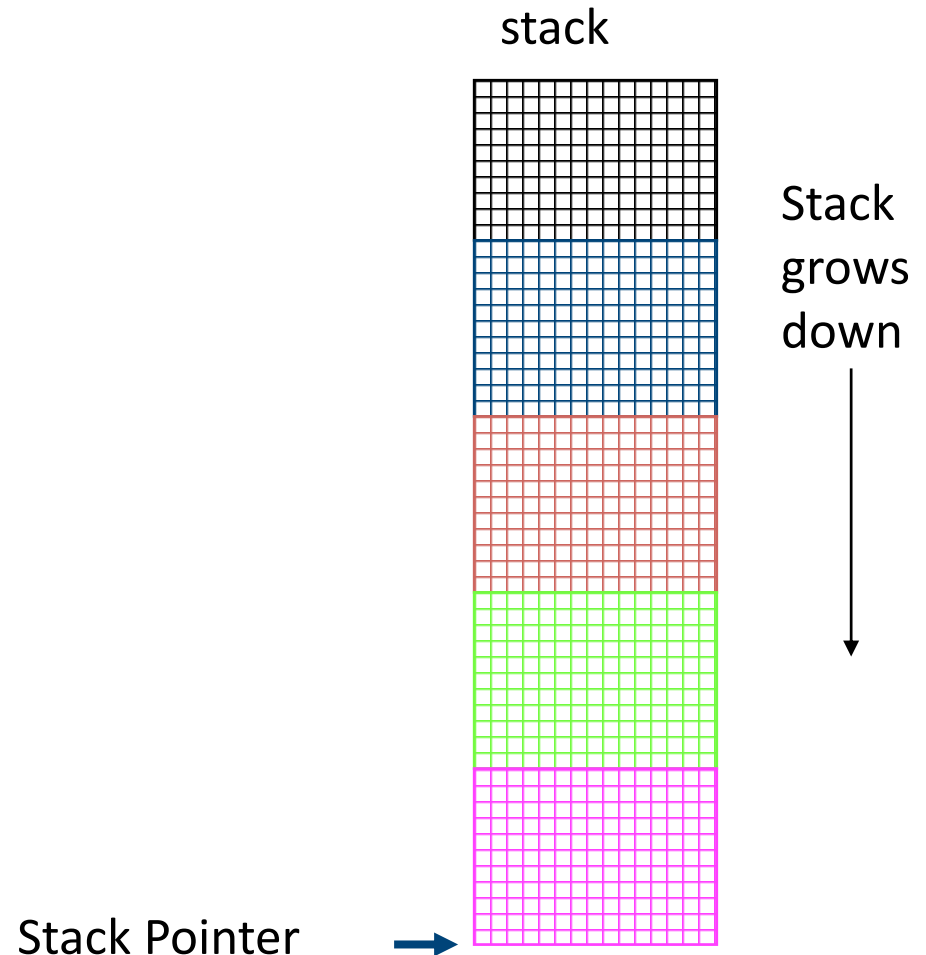
```
main () {  
    a(0);  
}  
void a (int m) {  
    b(1);  
}  
void b (int n) {  
    c(2);  
}  
void c (int o) {  
    d(3);  
}  
void d (int p) {  
}
```



Stack

- ▶ Last In, First Out (LIFO) data structure

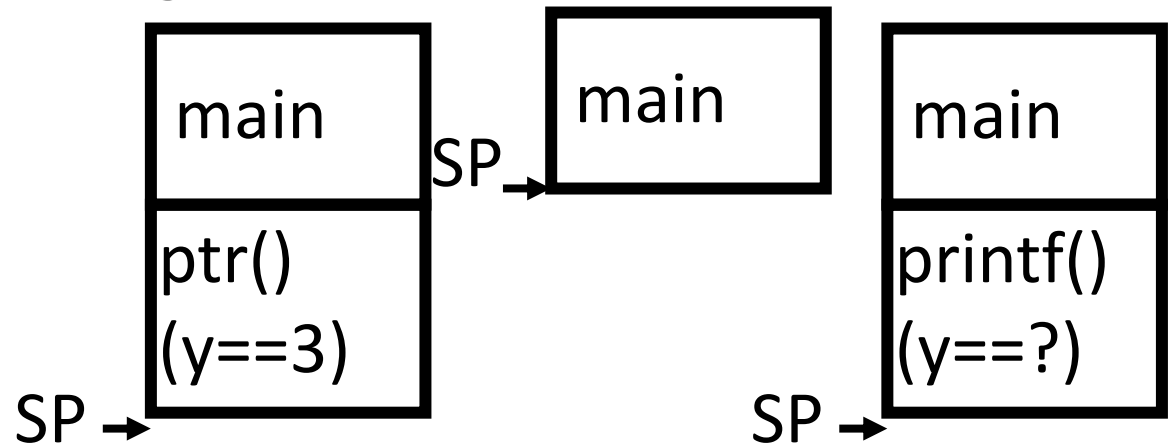
```
main () {  
    a(0);  
}  
void a (int m) {  
    b(1);  
}  
void b (int n) {  
    c(2);  
}  
void c (int o) {  
    d(3);  
}  
void d (int p) {  
}
```



Who cares about stack management?

- ▶ Pointers in C allow access to deallocated memory, leading to hard-to-find bugs !

```
int *ptr () {  
    int y;  
    y = 3;  
    return &y;  
};
```



```
main () {  
    int *stackAddr, content;  
    stackAddr = ptr();  
    content = *stackAddr;  
    printf("%d", content); /* 3 */  
    content = *stackAddr;  
    printf("%d", content); /* -2 */  
};
```


The Heap (Dynamic memory)

- ▶ Large pool of memory, not allocated in contiguous order
 - back-to-back requests for heap memory could result blocks very far apart
 - where Java/C++ `new` command allocates memory
- ▶ In C, specify number of bytes of memory explicitly to allocate item

```
int *ptr;  
ptr = (int *) malloc(sizeof(int));  
/* malloc returns type (void *),  
so need to cast to right type */
```

- `malloc()`: Allocates raw, uninitialized memory from heap

Memory Management

- ▶ How do we manage memory?
 - Code, Static
 - Simple
 - They never grow or shrink
 - Stack
 - Simple
 - Stack frames are created and destroyed in last-in, first-out (LIFO) order
 - Heap
 - Tricky
 - Memory can be allocated / deallocated at any time

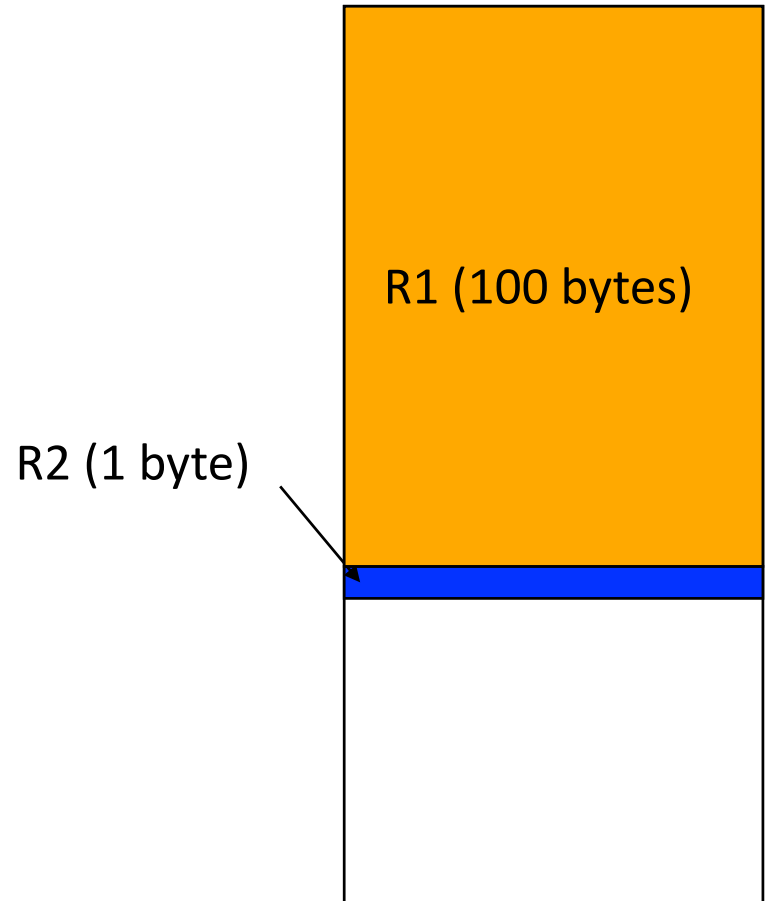
Heap Management Requirements

- ▶ Want `malloc()` and `free()` to run quickly.
- ▶ Want minimal memory overhead
- ▶ Want to avoid **fragmentation***
 - When most of our free memory is in many small chunks
 - In this case, we might have many free bytes but not be able to satisfy a large request since the free bytes are not contiguous in memory.

* This is technically called *external fragmentation*

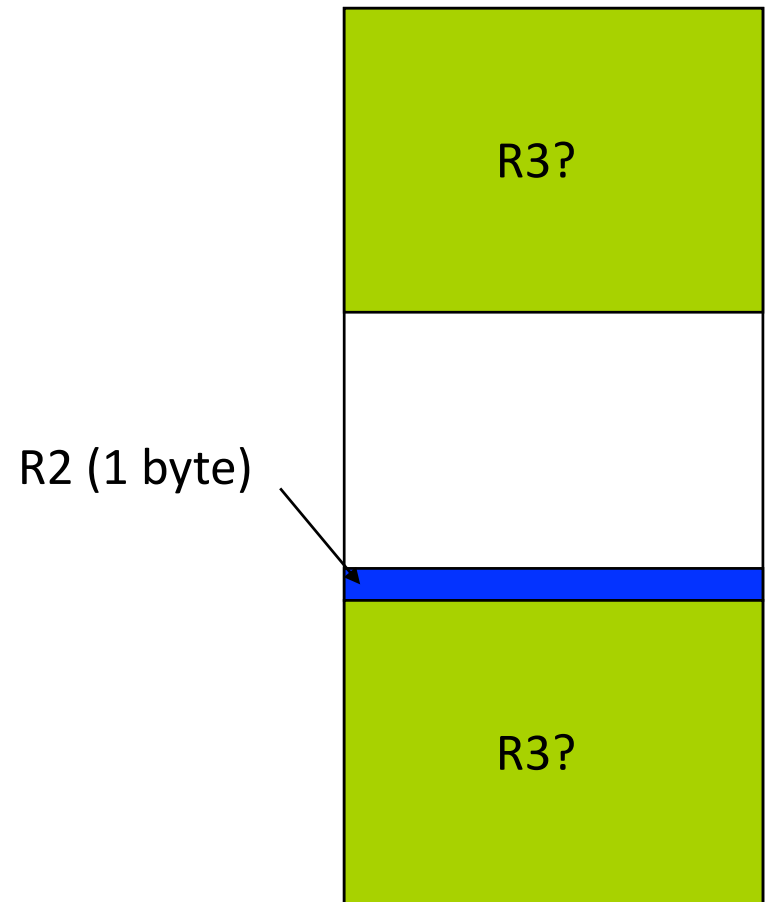
Heap Management

- ▶ An example
 - Request R1 for 100 bytes
 - Request R2 for 1 byte
 - Memory from R1 is freed



Heap Management

- ▶ An example
 - Request R1 for 100 bytes
 - Request R2 for 1 byte
 - Memory from R1 is freed
 - Request R3 for 50 bytes



K&R Malloc/Free Implementation

- ▶ From Section 8.7 of K&R
 - Code in the book uses some C language features we haven't discussed and is written in a very terse style, don't worry if you can't decipher the code
- ▶ Each block of memory is preceded by a header that has two fields:
 - **size** of the block
 - a **pointer to the next** block
- ▶ All **free blocks** are kept in a circular linked list, the pointer field is unused in an allocated block

K&R Implementation

- ▶ `malloc()` searches the free list for a block that is big enough. If none is found, more memory is requested from the operating system. If what it gets can't satisfy the request, it fails.
- ▶ `free()` checks if the blocks adjacent to the freed block are also free
 - If so, adjacent free blocks are merged (**coalesced**) into a single, larger free block
 - Otherwise, the freed block is just added to the free list

Choosing a block in `malloc()`

- ▶ If there are multiple free blocks of memory that are big enough for some request, how do we choose which one to use?
 - **best-fit**: choose the smallest block that is big enough for the request
 - **first-fit**: choose the first block we see that is big enough
 - **next-fit**: like first-fit but remember where we finished searching and resume searching from there

Tradeoffs of allocation policies

- ▶ **Best-fit:** Tries to limit fragmentation but at the cost of time (must examine all free blocks for each malloc).
 - Leaves lots of small blocks (why?)
- ▶ **First-fit:** Quicker than best-fit (why?) but potentially more fragmentation.
 - Tends to concentrate small blocks at the beginning of the free list (why?)
- ▶ **Next-fit:** Does not concentrate small blocks at front like first-fit, should be faster as a result.

Quiz – Pros and Cons of fits

- 1) **first-fit** results in many **small blocks** at the beginning of the free list
- 2) **next-fit** is **slower than first-fit**, since it takes longer in steady state to find a match
- 3) **best-fit** leaves lots of tiny blocks

	1	2	3
a)	F	F	T
b)	F	T	T
c)	T	F	F
d)	T	F	T
e)	T	T	T

Quiz – Pros and Cons of fits

- 1) **first-fit** results in many **small blocks** at the beginning of the free list
- 2) **next-fit** is **slower than first-fit**, since it takes longer in steady state to find a match
- 3) **best-fit** leaves lots of tiny blocks

- | | 1 | 2 | 3 |
|-----------|----------|----------|----------|
| a) | F | F | T |
| b) | F | T | T |
| c) | T | F | F |
| d) | T | F | T |
| e) | T | T | T |

Summary

- ▶ C has 3 pools of memory
 - Static storage: global variable storage, basically permanent, entire program run
 - The Stack: local variable storage, parameters, return address
 - The Heap (dynamic storage): `malloc()` grabs space from here, `free()` returns it.
- ▶ `malloc()` handles free space with freelist. Three different ways to find free space when given a request:
 - **First fit** (find first one that's free)
 - **Next fit** (same as first, but remembers where left off)
 - **Best fit** (finds most “snug” free space)