

CSE 31

Computer Organization

Lecture 12 – MIPS: Logical Operators
Machine Instruction Format (1)



Announcement

- ▶ Lab #6 this week
 - Due next week
- ▶ HW #3 (from zyBooks)
 - Due Monday (10/8)
- ▶ Project #1
 - Due Monday (10/22)
 - Don't start late, you won't have time!
- ▶ Reading assignment
 - Chapter 4.1-4.9 of zyBooks (Reading Assignment #4)
 - Make sure to do the Participation Activities
 - Due Monday (10/15)

Review

- ▶ Functions called with **jal**, return with **jr \$ra**.
- ▶ The stack is your friend: Use it to save anything you need. Just leave it the way you found it!
- ▶ Instructions we know so far...
 - Arithmetic: **add, addi, sub, addu, addiu, subu**
 - Memory: **lw, sw, lb, sb, lbu**
 - Decision: **beq, bne, slt, slti, sltu, sltiu**
 - Unconditional Branches (Jumps): **j, jal, jr**
- ▶ Registers we know so far
 - All of them!
 - There are CONVENTIONS when calling procedures!

Quiz

```
r:  ...      # R/W $s0,$v0,$t0,$a0,$sp,$ra,mem
    ...      ### PUSH REGISTER(S) TO STACK?
    jal e     # Call e
    ...      # R/W $s0,$v0,$t0,$a0,$sp,$ra,mem
    jr $ra    # Return to caller of r

e:  ...      # R/W $s0,$v0,$t0,$a0,$sp,$ra,mem
    jr $ra    # Return to r
```

What does `r` have to push on the stack before “`jal e`”?

- a) 1 of (\$s0, \$sp, \$v0, \$t0, \$a0, \$ra)
- b) 2 of (\$s0, \$sp, \$v0, \$t0, \$a0, \$ra)
- c) 3 of (\$s0, \$sp, \$v0, \$t0, \$a0, \$ra)
- d) 4 of (\$s0, \$sp, \$v0, \$t0, \$a0, \$ra)
- e) 5 of (\$s0, \$sp, \$v0, \$t0, \$a0, \$ra)

Quiz

```
r:  ...      # R/W $s0, $v0, $t0, $a0, $sp, $ra, mem
    ...      ### PUSH REGISTER(S) TO STACK?
    jal e     # Call e
    ...      # R/W $s0, $v0, $t0, $a0, $sp, $ra, mem
    jr $ra    # Return to caller of r

e:  ...      # R/W $s0, $v0, $t0, $a0, $sp, $ra, mem
    jr $ra    # Return to r
```

What does `r` have to push on the stack before “`jal e`”?

Saved Volatile! -- need to push

- a) 1 of (\$s0, \$sp, \$v0, \$t0, \$a0, \$ra)
- b) 2 of (\$s0, \$sp, \$v0, \$t0, \$a0, \$ra)
- c) 3 of (\$s0, \$sp, \$v0, \$t0, \$a0, \$ra)
- 😊 d) 4 of (\$s0, \$sp, \$v0, \$t0, \$a0, \$ra)
- e) 5 of (\$s0, \$sp, \$v0, \$t0, \$a0, \$ra)

Bitwise Operations

- ▶ So far, we've done arithmetic (**add**, **sub**, **addi**), mem access (**lw** and **sw**), & branches and jumps.
- ▶ All of these instructions view contents of register as a single quantity (e.g., signed or unsigned int)
- ▶ **New Perspective**: View register as 32 raw bits rather than as a single 32-bit number
 - Since registers are composed of 32 bits, wish to access individual bits (or groups of bits) rather than the whole.
- ▶ Introduce two new classes of instructions
 - **Logical & Shift Ops**

Logical Operators (1/3)

- ▶ Two basic logical operators:
 - AND: outputs 1 only if **all** inputs are 1
 - OR: outputs 1 if **at least one** input is 1
- ▶ Truth Table: standard table listing all possible combinations of inputs and resultant output

A	B	A AND B	A OR B
0	0	0	0
0	1	0	1
1	0	0	1
1	1	1	1

Logical Operators (2/3)

- ▶ Logical Instruction Syntax:
1 2,3,4
 - where
 - 1) operation name
 - 2) register that will receive value
 - 3) first operand (register)
 - 4) second operand (register) or immediate (numerical constant)
- ▶ In general, can define them to accept > 2 inputs, but in the case of MIPS assembly, these accept exactly 2 inputs and produce 1 output
 - Again, rigid syntax, simpler hardware

Logical Operators (3/3)

▶ Instruction Names:

- **and, or**: Both of these expect the third argument to be a register
- **andi, ori**: Both of these expect the third argument to be an immediate

▶ MIPS Logical Operators are all **bitwise**, meaning that bit 0 of the output is produced by the respective bit 0's of the inputs, bit 1 by the bit 1's, etc.

- C: Bitwise AND is `&` (e.g., **`z = x & y;`**)
- C: Bitwise OR is `|` (e.g., **`z = x | y;`**)

Uses of Logical Operators (1/3)

- ▶ Note that **and**ing a bit with 0 produces a 0 at the output while **and**ing a bit with 1 produces the original bit.
- ▶ This can be used to create a **mask**.

- Example:

1011	0110	1010	0100	0011	1101	1001	1010
0000	0000	0000	0000	0000	1111	1111	1111

- The result of **and**ing these:

0000	0000	0000	0000	0000	1101	1001	1010
------	------	------	------	------	------	------	------

mask:

1101	1001	1010
1111	1111	1111
1101	1001	1010

mask last 12 bits

Uses of Logical Operators (2/3)

- ▶ The second bitstring in the example is called a **mask**. It is used to isolate the rightmost 12 bits of the first bitstring by masking out the rest of the string (e.g. setting to all 0s).
- ▶ Thus, the **and** operator can be used to set certain portions of a bitstring to 0s, while leaving the rest alone.
 - In particular, if the first bitstring in the above example were in `$t0`, then the following instruction would mask it:
`andi $t0, $t0, 0xFFF`

Uses of Logical Operators (3/3)

- ▶ Similarly, note that **OR**ing a bit with **1** sets a **1** at the output while **OR**ing a bit with **0** keeps the original bit.
- ▶ Often used to force certain bits to **1**s.
 - For example, if `$t0` contains `0x12345678`, then after this instruction:

```
ori    $t0, $t0, 0xFFFF
```


... `$t0` will contain `0x1234FFFF`
 - (i.e., the high-order 16 bits are untouched, while the low-order 16 bits are forced to **1**s).

Shift Instructions (review) (1/4)

- ▶ Move (shift) all the bits in a word to the left or right by a number of bits.

- Example: shift right by 8 bits

1001 0010 0011 0100 0101 0110 0111 1000

0000 0000 1001 0010 0011 0100 0101 0110

- Example: shift left by 8 bits

0001 0010 1011 0100 0101 0110 0111 1000

1011 0100 0101 0110 0111 1000 0000 0000

Shift Instructions (2/4)

- ▶ Shift Instruction Syntax:

1 2,3,4

...where

1) operation name

2) register that will receive value

3) first operand (register)

4) shift amount (constant < 32)

- ▶ MIPS shift instructions:

1. **sll** (shift left logical): shifts left and fills emptied bits with 0s

2. **srl** (shift right logical): shifts right and fills emptied bits with 0s

3. **sra** (shift right arithmetic): shifts right and fills emptied bits by sign extending

Shift Instructions (3/4)

- ▶ Example: shift right arithmetic by 8 bits

 0001 0010 0011 0100 0101 0110 0111 1000

0000 0000 0001 0010 0011 0100 0101 0110

- ▶ Example: shift right arithmetic by 8 bits

 1001 0010 0011 0100 0101 0110 0111 1000

1111 1111 1001 0010 0011 0100 0101 0110

Shift Instructions (4/4)

- ▶ Since shifting is faster than multiplication, a good compiler usually notices when C code multiplies by a power of 2 and compiles it to a shift instruction:

`a *= 8;` (in C)

would compile to:

`sll $s0, $s0, 3` (in MIPS)

- ▶ Likewise, shift right to divide by powers of 2 (rounds towards $-\infty$)
 - remember to use `sra`

Summary

▶ Logical and Shift Instructions

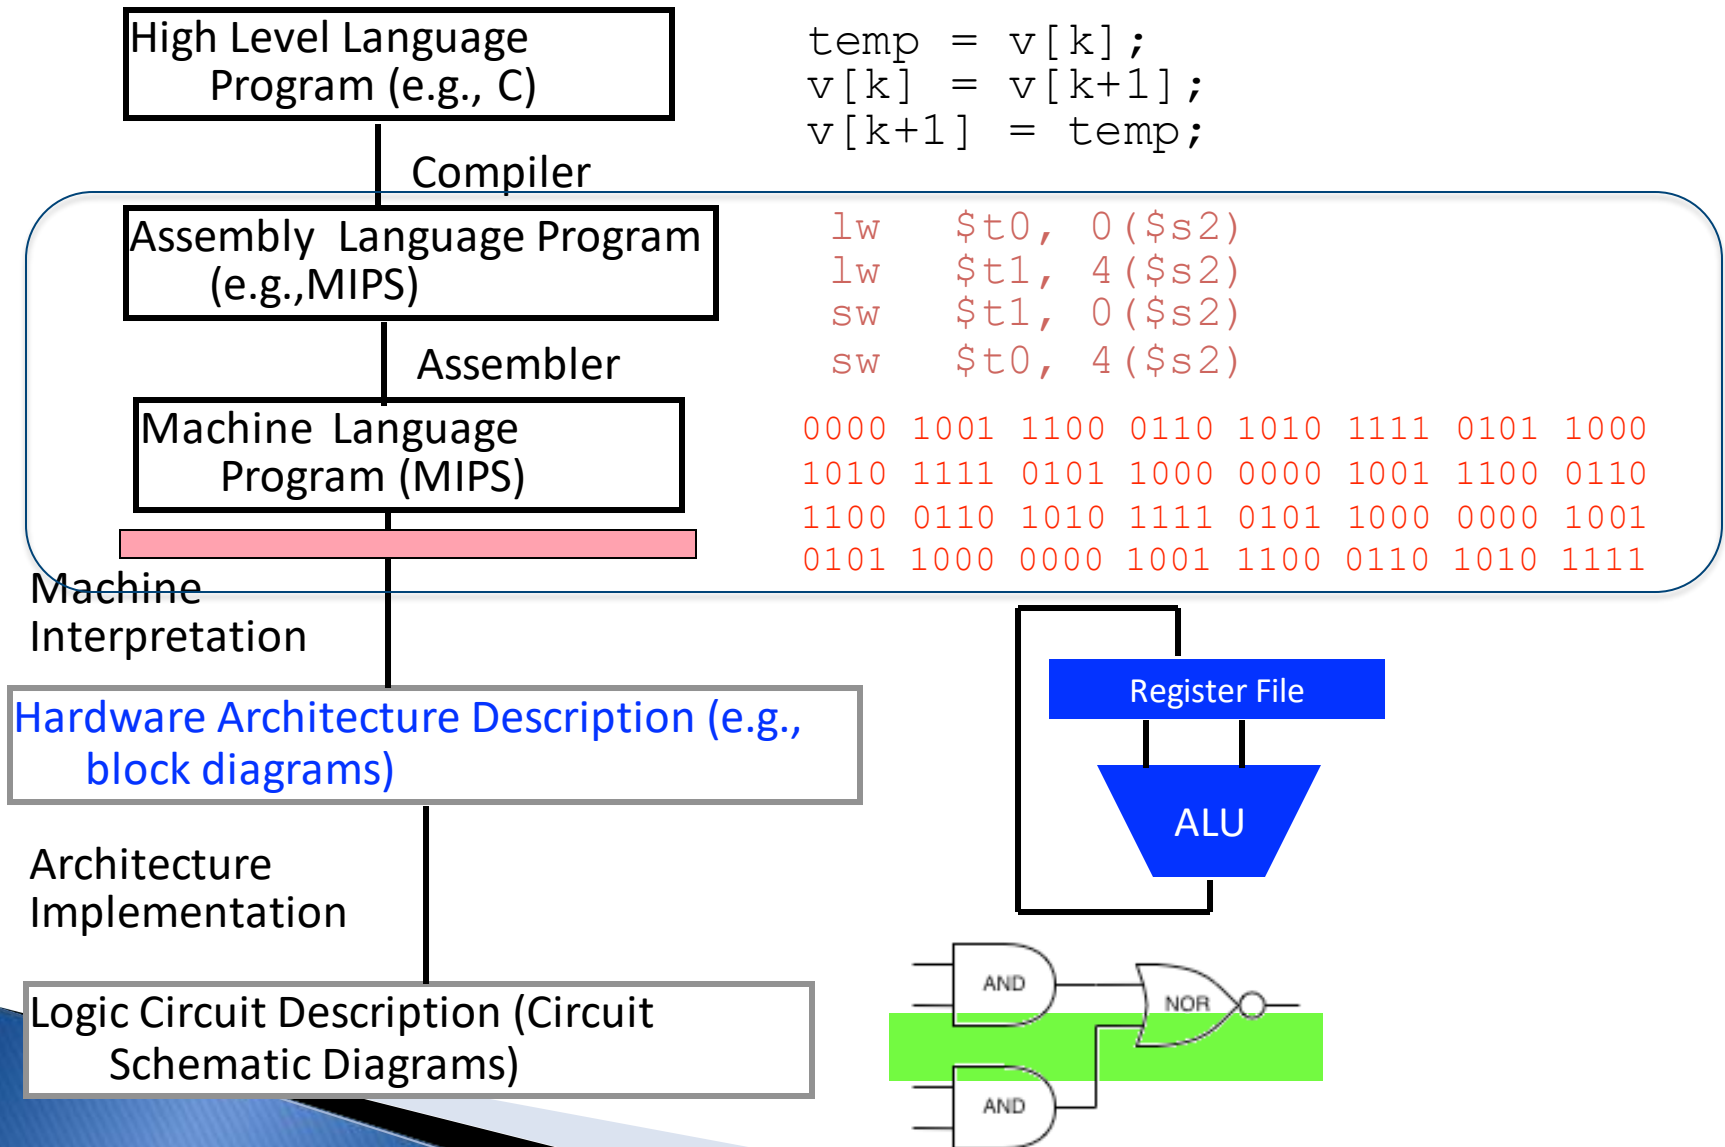
- Operate on bits individually, unlike arithmetic, which operate on entire word.
- Use to isolate fields, either by masking or by shifting back and forth.
- Use shift left logical, **sll**, for multiplication by powers of 2
- Use shift right logical, **srl**, for division by powers of 2 of unsigned numbers (**unsigned int**)
- Use shift right arithmetic, **sra**, for division by powers of 2 of signed numbers (**int**)

▶ New Instructions:

and, andi, or, ori, sll, srl, sra

▶ That's all you need to know about MIPS!

Levels of Representation (abstractions)



Overview – Instruction Representation

- ▶ Big idea: stored program
 - consequences of stored program
- ▶ Instructions as numbers
- ▶ Instruction encoding
- ▶ MIPS instruction format for arithmetic instructions
- ▶ MIPS instruction format for Immediate, Data transfer instructions

Big Idea: Stored-Program Concept

- ▶ Where are programs stored when they are being run?
 - How are they stored in memory?
- ▶ Computers built on 2 key principles:
 - Instructions are represented as bit patterns - can think of these as numbers.
 - Therefore, entire programs can be stored in memory to be read or written just like data.
- ▶ Simplifies SW/HW of computer systems:
 - Memory technology for data also used for programs

Consequence #1: Everything Addressed

- ▶ Since all instructions and data are stored in memory, everything has a memory address: instructions, data words
 - both branches and jumps use these
- ▶ C pointers are just memory addresses: they can point to anything in memory
 - Unconstrained use of addresses can lead to nasty bugs; up to you in C; limits in Java
- ▶ One register keeps address of instruction being executed: “Program Counter” (PC)
 - Basically a pointer to memory: Intel calls it Instruction Address Pointer, a better name

Consequence #2: Binary Compatibility

- ▶ Programs are distributed in binary form
 - Programs bound to specific instruction set
 - Different version for **Macintoshes** and **PCs**
- ▶ New machines want to run old programs (“binaries”) as well as programs compiled to new instructions
 - Leads to “backward compatible” instruction set evolving over time
 - Selection of Intel 8086 in 1981 for 1st IBM PC is major reason latest PCs still use 80x86 instruction set (Pentium 4); could still run program from 1981 PC today

Instructions as Numbers (1/2)

- ▶ Currently all data we work with is in words (32-bit blocks):
 - Each register is a word.
 - **lw** and **sw** both access memory one word at a time.
- ▶ So how do we represent instructions?
 - Remember: Computer only understands 1s and 0s, so “**add \$t0, \$0, \$0**” is meaningless.
 - MIPS wants simplicity: since data is in words, make instructions into words too!

Instructions as Numbers (2/2)

- ▶ One word is 32 bits, so divide instruction word into “fields”.
- ▶ Each field tells processor something about the instruction.
- ▶ We could define different fields for each instruction, but MIPS is based on simplicity, so define 3 basic types of instruction formats:
 - I-format
 - J-format
 - R-format
 - What do these letters (I, J, R) stand for?

Instruction Formats

- ▶ **I-format**: used for instructions with immediates, **lw** and **sw** (since offset counts as an immediate), and branches (**beq** and **bne**),
 - (but not the shift instructions; later)
- ▶ **J-format**: used for **j** and **jal**
- ▶ **R-format**: used for all other instructions
- ▶ Why 3 different formats?
 - It will soon become clear why the instructions have been partitioned in this way.

R-Format Instructions (1/5)

- ▶ Define “fields” of the following number of bits each: $6 + 5 + 5 + 5 + 5 + 6 = 32$

6	5	5	5	5	6
---	---	---	---	---	---

- ▶ For simplicity, each field has a name:

opcode	rs	rt	rd	shamt	funct
--------	----	----	----	-------	-------

- ▶ **Important:** On these slides and in book, each field is viewed as a 5- or 6-bit unsigned integer, not as part of a 32-bit integer.
 - Consequence: 5-bit fields can represent any number 0-31, while 6-bit fields can represent any number 0-63.

R-Format Instructions (2/5)

- ▶ What do these field integer values tell us?
 - `opcode`: partially specifies what instruction it is
 - Note: This number is equal to 0 for all R-Format instructions.
 - `funct`: combined with `opcode`, this number exactly specifies the instruction
- ▶ Question: Why aren't `opcode` and `funct` a single 12-bit field?
 - We'll answer this later.

R-Format Instructions (3/5)

- ▶ More fields:
 - rs (Source Register): **generally** used to specify register containing first operand
 - rt (Target Register): **generally** used to specify register containing second operand (note that name is misleading)
 - rd (Destination Register): **generally** used to specify register which will receive result of computation

R-Format Instructions (4/5)

- ▶ Notes about register fields:
 - Each register field is exactly 5 bits, which means that it can specify any unsigned integer in the range 0-31. Each of these fields specifies one of the 32 registers by number.
 - The word “generally” was used because there are exceptions
 - **mult** and **div** have nothing important in the **rd** field since the dest registers are **hi** and **lo**
 - **mfhi** and **mflo** have nothing important in the **rs** and **rt** fields since the source is determined by the instruction.

R-Format Instructions (5/5)

- ▶ Final field:
 - shamt: This field contains the amount a shift instruction will shift by. Shifting a 32-bit word by more than 31 is useless, so this field is only 5 bits (so it can represent the numbers 0-31).
 - This field is set to 0 in all but the shift instructions.
- ▶ For a detailed description of field usage for each instruction, see MIPS_Reference_Sheet.pdf in CatCourses (You can bring with you to all exams)

R-Format Example (1/2)

- ▶ MIPS Instruction:

add **\$t0, \$t1, \$t2**

add **\$8, \$9, \$10**

opcode = 0 (look up in table)

funct = 32 (look up in table)

rd = 8 (destination)

rs = 9 (first **operand**)

rt = 10 (second **operand**)

shamt = 0 (not a shift)

R-Format Example (2/2)

► MIPS Instruction:

add \$8, \$9, \$10

Decimal number per field representation:

0	9	10	8	0	32
---	---	----	---	---	----

Binary number per field representation:

000000	01001	01010	01000	00000	100000
--------	-------	-------	-------	-------	--------

hex

hex representation: 012A 4020_{hex}

decimal representation: 19,546,144_{ten}

Called a Machine Language Instruction

I-Format Instructions (1/4)

- ▶ What about instructions with immediates?
 - 5-bit field only represents numbers up to the value 31: immediates may be much larger than this
 - Ideally, MIPS would have only one instruction format (for simplicity): unfortunately, we need to compromise
- ▶ Define new instruction format that is partially consistent with R-format:
 - First notice that, if instruction has immediate, then it uses at most 2 registers.

I-Format Instructions (2/4)

- ▶ Define “fields” of the following number of bits each: $6 + 5 + 5 + 16 = 32$ bits

6	5	5	16
---	---	---	----

- Again, each field has a name:

opcode	rs	rt	immediate
--------	----	----	-----------

- **Key Concept:** Only one field is inconsistent with R-format. Most importantly, `opcode` is still in same location.

I-Format Instructions (3/4)

- ▶ What do these fields mean?
 - opcode: same as before except that, since there's no `funct` field, `opcode` uniquely specifies an instruction in I-format
 - This also answers question of why R-format has two 6-bit fields to identify instruction instead of a single 12-bit field: in order to be consistent as possible with other formats while leaving as much space as possible for immediate field.
 - rs: specifies a register operand (if there is one)
 - rt: specifies register which will receive result of computation (this is why it's called the **target** register “`rt`”) or other operand for some instructions.

I-Format Instructions (4/4)

▶ The Immediate Field:

- **addi, slti, sltiu**, the immediate is **sign-extended** to 32 bits. Thus, it's treated as a signed integer.
- 16 bits → can be used to represent immediate up to 2^{16} different values
- This is large enough to handle the offset in a typical **lw** or **sw**, plus a vast majority of values that will be used in the **slti** instruction.
- We'll see what to do when the number is too big later

I-Format Example (1/2)

- ▶ MIPS Instruction:

addi **\$s5 , \$s6 , -50**

addi **\$21 , \$22 , -50**

opcode = 8 (look up in table in book)

rs = 22 (register containing operand)

rt = 21 (target register)

immediate = -50 (by default, this is decimal)

I-Format Example (2/2)

- ▶ MIPS Instruction:

`addi $21, $22, -50`

Decimal/field representation:

8	22	21	-50
---	----	----	-----

Binary/field representation:

001000	10110	10101	1111111111001110
--------	-------	-------	------------------

hexadecimal representation: `22D5 FFCEhex`

decimal representation: `584, 449, 998ten`

Quiz

Which instruction has same representation as 35_{ten} ?

- a) add \$0, \$0, \$0
- b) subu \$s0,\$s0,\$s0
- c) lw \$0, 0(\$0)
- d) addi \$0, \$0, 35
- e) subu \$0, \$0, \$0

opcode	rs	rt	rd	shamt	funct
opcode	rs	rt	rd	shamt	funct
opcode	rs	rt	offset		
opcode	rs	rt	immediate		
opcode	rs	rt	rd	shamt	funct

Registers numbers and names:

0: \$0, .. 8: \$t0, 9: \$t1, ..15: \$t7, 16: \$s0, 17: \$s1, .. 23: \$s7

Opcodes and function fields (if necessary)

add: opcode = 0, funct = 32

subu: opcode = 0, funct = 35

addi: opcode = 8

lw: opcode = 35

Quiz

Which instruction has same representation as 35_{ten} ?

- a) add \$0, \$0, \$0
- b) subu \$s0,\$s0,\$s0
- c) lw \$0, 0(\$0)
- d) addi \$0, \$0, 35
- e) subu \$0, \$0, \$0**

opcode	rs	rt	rd	shamt	funct
opcode	rs	rt	rd	shamt	funct
opcode	rs	rt	offset		
opcode	rs	rt	immediate		
opcode	rs	rt	rd	shamt	funct

Registers numbers and names:

0: \$0, .. 8: \$t0, 9: \$t1, ..15: \$t7, 16: \$s0, 17: \$s1, .. 23: \$s7

Opcodes and function fields (if necessary)

add: opcode = 0, funct = 32

subu: opcode = 0, funct = 35

addi: opcode = 8

lw: opcode = 35